

Cntxt

The AI Context Filter



Naveh Talor | Liav Samiya | Gal Shitrit | Daniel Armoni

The Context **Exposure** Problem


AI agents need access to enterprise knowledge, but unrestricted context sharing exposes confidential information.

- Enterprise data contains hidden secrets, PII, and financial metrics.
- Sending raw data to external models breaks compliance boundaries.
- Current architecture forces a choice between AI capability and data security.



The MCP Ecosystem Future

AI agents are rapidly connecting to internal systems – Slack, GitHub, Drive, and Databases – through the Model Context Protocol (MCP).



The **Missing Layer** in AI Security

Authentication connects the pipes, but it doesn't filter the water.



Access Control

"Who can access enterprise data?"

Solved by identity providers and API gateways.




Context Control

"What should the model actually see?"

The missing layer. Passing raw data exposes secrets, PII, and IP.

Complementary **Layers**, Not Competitors


Cntxt sits above traditional MCP gateways, focusing on data sanitization & context preservation.




The Cntxt Middleware

An AI-Safe Context Firewall

- Sits directly between your MCP connectors and external AI models.
- Intercepts raw enterprise knowledge in real-time.
- Transforms unsafe variables into anonymized placeholders.
- Ensures external models only compute on safe, abstract representations.



How It Works



Cntxt operates entirely in the middle (Stages 2-4), securing context before it ever leaves your trusted perimeter.

Context Transformation

Models need relationships and structure to reason, not the actual raw secrets.

RAW ENTERPRISE CONTEXT

"Please review the contract for **Acme Corp**. The total budget allocated is **\$4,500,000**. Contact **Sarah Jenkins (s.jenkins@acme.com)**. Use API key **sk-live-8x9b2v** for the deployment."

SAFE AI CONTEXT (CNTXT OUTPUT)

"Please review the contract for **COMPANY_1**. The total budget allocated is **AMOUNT_1**. Contact **PERSON_1 (EMAIL_1)**. Use API key **API_KEY_1** for the deployment."

Persistent Project Memory

Context across conversations, securely.

Simple redaction breaks AI reasoning because context is lost between separate prompts.

Cntxt innovation maintains a **stable mapping**. If "Sarah Jenkins" is PERSON_1 in document A, she remains PERSON_1 in document B within the same MCP project.

The model understands the entity relationship, but never learns the entity's true identity.

PERSISTENT MAPPING MECHANISM


DOCUMENT A "Sarah
Jenkins" **Mapped to**
PERSON_1

DOCUMENT B "Sarah
Jenkins" **Mapped to**
PERSON_1

PROJECT MEMORY
PERSON_1
ID remains stable across
sessions

Result: The AI assistant maintains perfect context and coherence when referencing "PERSON_1" across multiple separate prompts, without ever learning her actual name.

System Architecture



Live Demo

The Context Security Layer

Empowering the future of enterprise AI agents with uncompromising data security and privacy.

Cntxt Hackathon Team

OpenAI Build Week